

DAV_Python

April 24, 2024

1 Simple Linear Regression

```
[1]: # Essential imports
import numpy as np
import matplotlib.pyplot as plt
from sklearn import datasets
from sklearn.model_selection import train_test_split
import pandas as pd
import seaborn as sns
```

```
[2]: # Iris dataset
iris = sns.load_dataset('iris')
iris
```

```
[2]:
```

	sepal_length	sepal_width	petal_length	petal_width	species
0	5.1	3.5	1.4	0.2	setosa
1	4.9	3.0	1.4	0.2	setosa
2	4.7	3.2	1.3	0.2	setosa
3	4.6	3.1	1.5	0.2	setosa
4	5.0	3.6	1.4	0.2	setosa
..
145	6.7	3.0	5.2	2.3	virginica
146	6.3	2.5	5.0	1.9	virginica
147	6.5	3.0	5.2	2.0	virginica
148	6.2	3.4	5.4	2.3	virginica
149	5.9	3.0	5.1	1.8	virginica

[150 rows x 5 columns]


```
[3]: # Performing Simple LR on petal_length and petal_width
# Extract features
X = iris['petal_length']
Y = iris['petal_width']
X, Y
```

```
[3]: (0      1.4
      1      1.4
      2      1.3
```

```
3      1.5
4      1.4
...
145    5.2
146    5.0
147    5.2
148    5.4
149    5.1
Name: petal_length, Length: 150, dtype: float64,
0      0.2
1      0.2
2      0.2
3      0.2
4      0.2
...
145    2.3
146    1.9
147    2.0
148    2.3
149    1.8
Name: petal_width, Length: 150, dtype: float64)
```

```
[4]: # Visualise data
plt.scatter(X, Y)
```

```
[4]: <matplotlib.collections.PathCollection at 0x1ef42191d90>
```



```
[5]: # Split into training and testing sets
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, random_state=42,
↪test_size=0.1)
```

```
[6]: # Print training set
X_train, Y_train
```

```
[6]: (56      4.7
      104     5.8
       69     3.9
       55     4.5
      132     5.6
      ...
       71     4.0
      106     4.5
       14     1.2
       92     4.0
       102     5.9
      Name: petal_length, Length: 135, dtype: float64,
      56      1.6
      104     2.2
       69     1.1
```

```

55      1.3
132     2.2
...
71      1.3
106     1.7
14      0.2
92      1.2
102     2.1
Name: petal_width, Length: 135, dtype: float64)

```

```

[7]: # Print testing set
X_test, Y_test

```

```

[7]: (73      4.7
18      1.7
118     6.9
78      4.5
76      4.8
31      1.5
64      3.6
141     5.1
68      4.5
82      3.9
110     5.1
12      1.4
36      1.3
9       1.5
19      1.5
Name: petal_length, dtype: float64,
73      1.2
18      0.3
118     2.3
78      1.5
76      1.4
31      0.4
64      1.3
141     2.3
68      1.5
82      1.2
110     2.0
12      0.1
36      0.2
9       0.1
19      0.3
Name: petal_width, dtype: float64)

```

```
[8]: # Import Linear Regression model
from sklearn.linear_model import LinearRegression
slr_model = LinearRegression()
```

```
[9]: # Reshape inputs
X_train, X_test = np.array(X_train).reshape(-1, 1), np.array(X_test).
    ↪ reshape(-1, 1)

# Train model
slr_model.fit(X_train, Y_train)
```

```
[9]: LinearRegression()
```

```
[10]: # Make predictions
Y_pred = slr_model.predict(X_test)
```


```
[11]: # Evaluate model
from sklearn.metrics import mean_squared_error, mean_absolute_error

print(f"Mean Squared Error: {mean_squared_error(Y_test, Y_pred)}")
print(f"Mean Absolute Error: {mean_absolute_error(Y_test, Y_pred)}")
```

Mean Squared Error: 0.046489206286248065

Mean Absolute Error: 0.15873640858391244

```
[12]: # Final predictions
plt.scatter(X_train, Y_train, color='red')
plt.plot(X_test, Y_pred, color='black')
plt.show();
```



2 Multiple Linear Regression

```
[13]: # Essential imports
import numpy as np
import matplotlib.pyplot as plt
from sklearn import datasets
from sklearn.model_selection import train_test_split
import pandas as pd
import seaborn as sns
```

```
[14]: # Iris dataset
iris = sns.load_dataset('iris')
iris
```

```
[14]:
```

	sepal_length	sepal_width	petal_length	petal_width	species
0	5.1	3.5	1.4	0.2	setosa
1	4.9	3.0	1.4	0.2	setosa
2	4.7	3.2	1.3	0.2	setosa
3	4.6	3.1	1.5	0.2	setosa
4	5.0	3.6	1.4	0.2	setosa
..

145	6.7	3.0	5.2	2.3	virginica
146	6.3	2.5	5.0	1.9	virginica
147	6.5	3.0	5.2	2.0	virginica
148	6.2	3.4	5.4	2.3	virginica
149	5.9	3.0	5.1	1.8	virginica

[150 rows x 5 columns]

```
[15]: # For example, predict petal_width using sepal_length, sepal_width and
      ↪ petal_length
```

```
# Drop species
iris.drop('species', axis=1, inplace=True)
```

```
[16]: # Set features
X = iris[['sepal_length', 'sepal_width', 'petal_length']]
Y = iris['petal_width']
X, Y
```

```
[16]: (   sepal_length  sepal_width  petal_length
0           5.1           3.5           1.4
1           4.9           3.0           1.4
2           4.7           3.2           1.3
3           4.6           3.1           1.5
4           5.0           3.6           1.4
..          ...           ...           ...
145          6.7           3.0           5.2
146          6.3           2.5           5.0
147          6.5           3.0           5.2
148          6.2           3.4           5.4
149          5.9           3.0           5.1
```

[150 rows x 3 columns],

```
0    0.2
1    0.2
2    0.2
3    0.2
4    0.2
```

```
...
145  2.3
146  1.9
147  2.0
148  2.3
149  1.8
```

Name: petal_width, Length: 150, dtype: float64)

```
[17]: # Training and testing splits
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, random_state=42,
↳test_size=0.2)
```

```
[18]: # Training data
X_train, Y_train
```

```
[18]: (      sepal_length  sepal_width  petal_length
      22          4.6          3.6          1.0
      15          5.7          4.4          1.5
      65          6.7          3.1          4.4
      11          4.8          3.4          1.6
      42          4.4          3.2          1.3
      ..          ...          ...          ...
      71          6.1          2.8          4.0
     106          4.9          2.5          4.5
      14          5.8          4.0          1.2
      92          5.8          2.6          4.0
     102          7.1          3.0          5.9

      [120 rows x 3 columns],
      22      0.2
      15      0.4
      65      1.4
      11      0.2
      42      0.2
      ...
      71      1.3
     106      1.7
      14      0.2
      92      1.2
     102      2.1
      Name: petal_width, Length: 120, dtype: float64)
```

```
[19]: # Testing data
X_test, Y_test
```

```
[19]: (      sepal_length  sepal_width  petal_length
      73          6.1          2.8          4.7
      18          5.7          3.8          1.7
     118          7.7          2.6          6.9
      78          6.0          2.9          4.5
      76          6.8          2.8          4.8
      31          5.4          3.4          1.5
      64          5.6          2.9          3.6
     141          6.9          3.1          5.1
      68          6.2          2.2          4.5
```


82		5.8	2.7	3.9
110		6.5	3.2	5.1
12		4.8	3.0	1.4
36		5.5	3.5	1.3
9		4.9	3.1	1.5
19		5.1	3.8	1.5
56		6.3	3.3	4.7
104		6.5	3.0	5.8
69		5.6	2.5	3.9
55		5.7	2.8	4.5
132		6.4	2.8	5.6
29		4.7	3.2	1.6
127		6.1	3.0	4.9
26		5.0	3.4	1.6
128		6.4	2.8	5.6
131		7.9	3.8	6.4
145		6.7	3.0	5.2
108		6.7	2.5	5.8
143		6.8	3.2	5.9
45		4.8	3.0	1.4
30		4.8	3.1	1.6,
73	1.2			
18	0.3			
118	2.3			
78	1.5			
76	1.4			
31	0.4			
64	1.3			
141	2.3			
68	1.5			
82	1.2			
110	2.0			
12	0.1			
36	0.2			
9	0.1			
19	0.3			
56	1.6			
104	2.2			
69	1.1			
55	1.3			
132	2.2			
29	0.2			
127	1.8			
26	0.4			
128	2.1			
131	2.0			
145	2.3			

```
108     1.8
143     2.3
45      0.3
30      0.2
Name: petal_width, dtype: float64)
```

```
[20]: # Import Linear Regression model
from sklearn.linear_model import LinearRegression

mlr = LinearRegression()
```

```
[21]: # Train model
mlr.fit(X_train, Y_train)
```

```
[21]: LinearRegression()
```

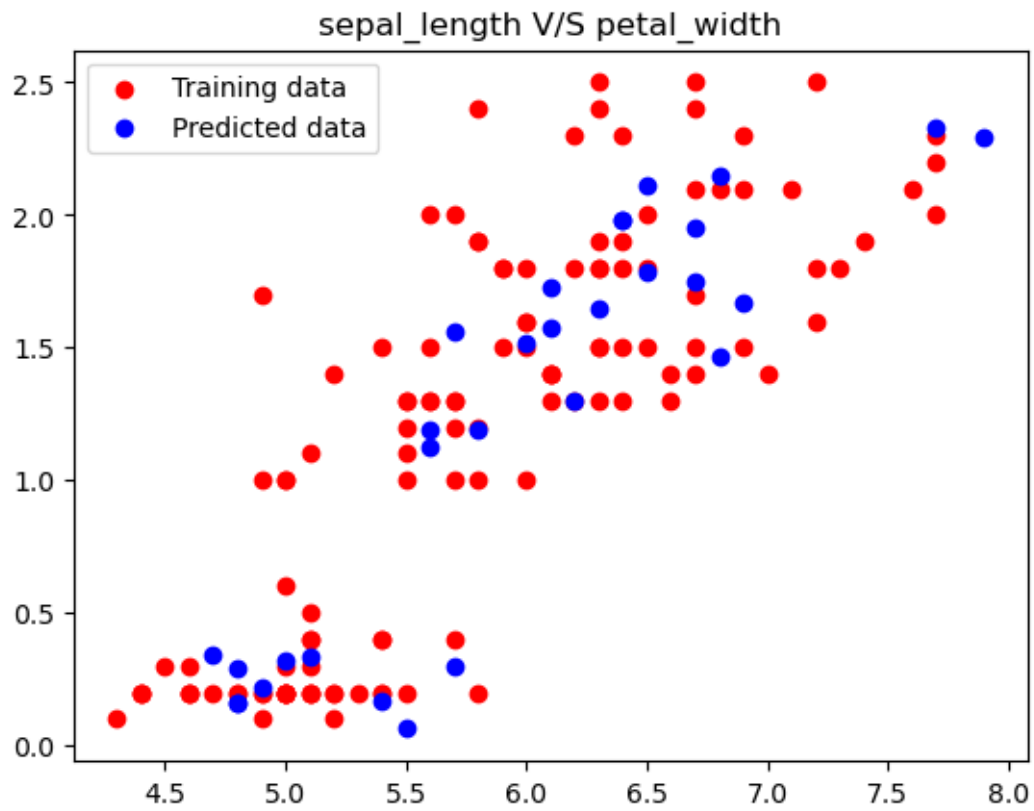
```
[22]: # Make predictions
Y_pred = mlr.predict(X_test)
```

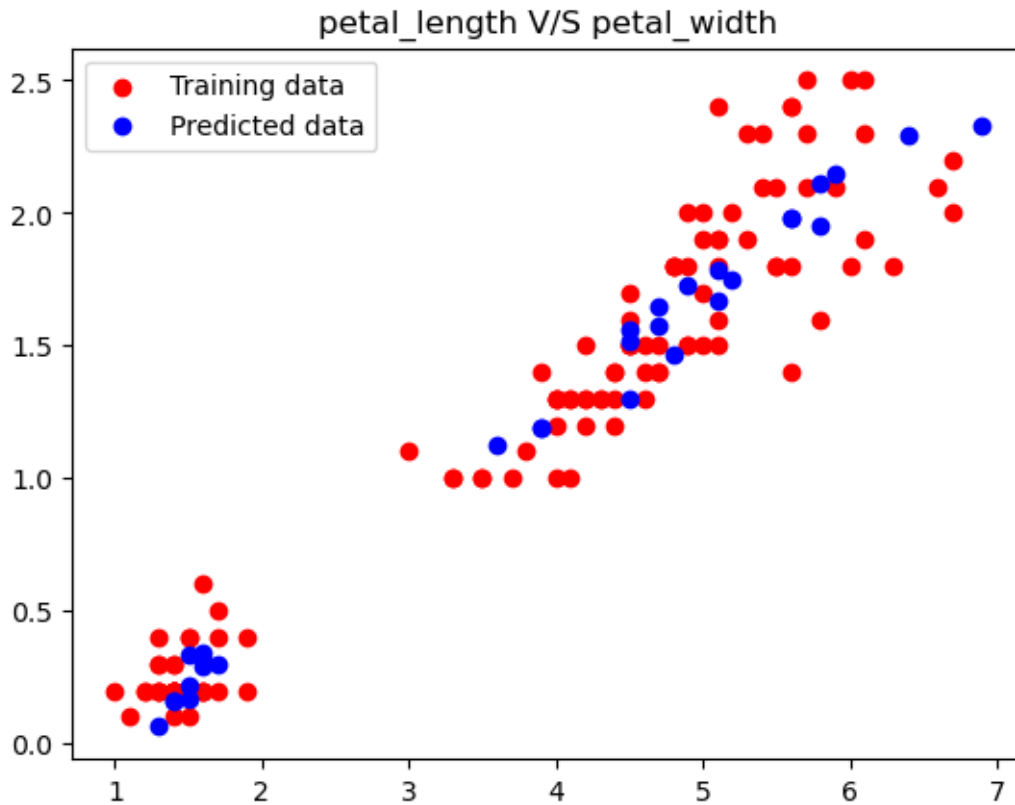
```
[23]: # Evaluate model
from sklearn.metrics import mean_squared_error, mean_absolute_error

print(f"Mean Squared Error: {mean_squared_error(Y_test, Y_pred)}")
print(f"Mean Absolute Error: {mean_absolute_error(Y_test, Y_pred)}")
```

```
Mean Squared Error: 0.046332603325764436
Mean Absolute Error: 0.15905558437496845
```

```
[24]: # Visualisations
for input_feature in X.columns:
    plt.scatter(X_train[input_feature], Y_train, color='red', label="Training_
↳data")
    plt.scatter(X_test[input_feature], Y_pred, color='blue', label="Predicted_
↳data")
    plt.title(f"{input_feature} V/S {Y.name}")
    plt.legend(loc="upper left")
    plt.show();
```





3 Logistic Regression

```
[25]: # Essential imports
import numpy as np
import matplotlib.pyplot as plt
from sklearn import datasets
from sklearn.model_selection import train_test_split
import pandas as pd
import seaborn as sns
```

```
[26]: # Penguins dataset
penguins = sns.load_dataset('penguins')
penguins
```

```
[26]:   species    island  bill_length_mm  bill_depth_mm  flipper_length_mm  \
0  Adelie  Torgersen         39.1          18.7           181.0
1  Adelie  Torgersen         39.5          17.4           186.0
2  Adelie  Torgersen         40.3          18.0           195.0
3  Adelie  Torgersen          NaN           NaN            NaN
4  Adelie  Torgersen         36.7          19.3           193.0
```

```

..      ...      ...      ...      ...      ...
339  Gentoo      Biscoe      NaN      NaN      NaN
340  Gentoo      Biscoe      46.8      14.3      215.0
341  Gentoo      Biscoe      50.4      15.7      222.0
342  Gentoo      Biscoe      45.2      14.8      212.0
343  Gentoo      Biscoe      49.9      16.1      213.0

```

```

      body_mass_g      sex
0      3750.0      Male
1      3800.0      Female
2      3250.0      Female
3      NaN      NaN
4      3450.0      Female
..      ...      ...
339      NaN      NaN
340      4850.0      Female
341      5750.0      Male
342      5200.0      Female
343      5400.0      Male

```

[344 rows x 7 columns]

```

[27]: # Remove NaN values
penguins.dropna(inplace=True)
penguins

```

```

[27]:      species      island  bill_length_mm  bill_depth_mm  flipper_length_mm  \
0      Adelie  Torgersen      39.1      18.7      181.0
1      Adelie  Torgersen      39.5      17.4      186.0
2      Adelie  Torgersen      40.3      18.0      195.0
4      Adelie  Torgersen      36.7      19.3      193.0
5      Adelie  Torgersen      39.3      20.6      190.0
..      ...      ...
338  Gentoo      Biscoe      47.2      13.7      214.0
340  Gentoo      Biscoe      46.8      14.3      215.0
341  Gentoo      Biscoe      50.4      15.7      222.0
342  Gentoo      Biscoe      45.2      14.8      212.0
343  Gentoo      Biscoe      49.9      16.1      213.0

      body_mass_g      sex
0      3750.0      Male
1      3800.0      Female
2      3250.0      Female
4      3450.0      Female
5      3650.0      Male
..      ...      ...
338      4925.0      Female

```

```

340      4850.0  Female
341      5750.0   Male
342      5200.0  Female
343      5400.0   Male

```

[333 rows x 7 columns]

```

[28]: # Predict sex using bill_length_mm, bill_depth_mm, flipper_length_mm and
      ↪ body_mass_g

      # Encode sex
      penguins.replace("Male", 0, inplace=True)
      penguins.replace("Female", 1, inplace=True)
      penguins

```

```

[28]:   species      island  bill_length_mm  bill_depth_mm  flipper_length_mm  \
0   Adelie  Torgersen         39.1          18.7           181.0
1   Adelie  Torgersen         39.5          17.4           186.0
2   Adelie  Torgersen         40.3          18.0           195.0
4   Adelie  Torgersen         36.7          19.3           193.0
5   Adelie  Torgersen         39.3          20.6           190.0
..      ...      ...
338  Gentoo    Biscoe         47.2          13.7           214.0
340  Gentoo    Biscoe         46.8          14.3           215.0
341  Gentoo    Biscoe         50.4          15.7           222.0
342  Gentoo    Biscoe         45.2          14.8           212.0
343  Gentoo    Biscoe         49.9          16.1           213.0

      body_mass_g  sex
0          3750.0    0
1          3800.0    1
2          3250.0    1
4          3450.0    1
5          3650.0    0
..      ...
338         4925.0    1
340         4850.0    1
341         5750.0    0
342         5200.0    1
343         5400.0    0

```

[333 rows x 7 columns]

```

[29]: # Features
      X = penguins[['bill_length_mm', 'bill_depth_mm', 'flipper_length_mm',
      ↪ 'body_mass_g']]
      Y = penguins['sex']

```

```
X, Y
```

```
[29]: (      bill_length_mm  bill_depth_mm  flipper_length_mm  body_mass_g
0           39.1           18.7           181.0          3750.0
1           39.5           17.4           186.0          3800.0
2           40.3           18.0           195.0          3250.0
4           36.7           19.3           193.0          3450.0
5           39.3           20.6           190.0          3650.0
..          ...          ...          ...          ...
338          47.2           13.7           214.0          4925.0
340          46.8           14.3           215.0          4850.0
341          50.4           15.7           222.0          5750.0
342          45.2           14.8           212.0          5200.0
343          49.9           16.1           213.0          5400.0
```

```
[333 rows x 4 columns],
```

```
0      0
```

```
1      1
```

```
2      1
```

```
4      1
```

```
5      0
```

```
..
```

```
338     1
```

```
340     1
```

```
341     0
```

```
342     1
```

```
343     0
```

```
Name: sex, Length: 333, dtype: int64)
```

```
[30]: # Training and testing split
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, random_state=42,
↪test_size=0.1)
```

```
[31]: # Training set
X_train, Y_train
```

```
[31]: (      bill_length_mm  bill_depth_mm  flipper_length_mm  body_mass_g
285           49.8           16.8           230.0          5700.0
296           47.5           14.2           209.0          4600.0
188           47.6           18.3           195.0          3850.0
260           42.7           13.7           208.0          3950.0
52           35.0           17.9           190.0          3450.0
..          ...          ...          ...          ...
194           50.9           19.1           196.0          3550.0
77           37.2           19.4           184.0          3900.0
112           39.7           17.7           193.0          3200.0
277           45.5           15.0           220.0          5000.0
```


108	38.1	17.0	181.0	3175.0
-----	------	------	-------	--------

[299 rows x 4 columns],

285 0

296 1

188 1

260 1

52 1

..

194 0

77 0

112 1

277 0

108 1

Name: sex, Length: 299, dtype: int64)

[32]: # Testing set

X_test, Y_test

[32]:	(bill_length_mm	bill_depth_mm	flipper_length_mm	body_mass_g
	30	39.5	16.7	178.0	3250.0
	317	46.9	14.6	222.0	4875.0
	79	42.1	19.1	195.0	4000.0
	201	49.8	17.3	198.0	3675.0
	63	41.1	18.2	192.0	4050.0
	304	44.9	13.8	212.0	4750.0
	289	50.7	15.0	223.0	5550.0
	186	49.7	18.6	195.0	3600.0
	217	49.6	18.2	193.0	3775.0
	203	51.4	19.0	201.0	3950.0
	81	42.9	17.6	196.0	4700.0
	14	34.6	21.1	198.0	4400.0
	328	43.3	14.0	208.0	4575.0
	132	36.8	18.5	193.0	3500.0
	272	45.1	14.4	210.0	4400.0
	138	37.0	16.5	185.0	3400.0
	120	36.2	17.2	187.0	3150.0
	152	46.5	17.9	192.0	3500.0
	82	36.7	18.8	187.0	3800.0
	282	45.7	13.9	214.0	4400.0
	115	42.7	18.3	196.0	4075.0
	143	40.7	17.0	190.0	3725.0
	323	49.1	15.0	228.0	5500.0
	205	50.7	19.7	203.0	4050.0
	6	38.9	17.8	181.0	3625.0
	116	38.6	17.0	188.0	2900.0
	268	44.9	13.3	213.0	5100.0

332		43.5	15.2	213.0	4650.0
169		58.0	17.8	181.0	3700.0
331		49.8	15.9	229.0	5950.0
174		43.2	16.6	187.0	2900.0
309		52.1	17.0	230.0	5550.0
69		41.8	19.4	198.0	4450.0
90		35.7	18.0	202.0	3550.0,
30	1				
317	1				
79	0				
201	1				
63	0				
304	1				
289	0				
186	0				
217	0				
203	0				
81	0				
14	0				
328	1				
132	1				
272	1				
138	1				
120	1				
152	1				
82	1				
282	1				
115	0				
143	0				
323	0				
205	0				
6	1				
116	1				
268	1				
332	1				
169	1				
331	0				
174	1				
309	0				
69	0				
90	1				

Name: sex, dtype: int64)

```
[33]: # Import Logistic Regression model
from sklearn.linear_model import LogisticRegression

lr_model = LogisticRegression()
```

```
[34]: # Train model
lr_model.fit(X_train, Y_train)
```

```
[34]: LogisticRegression()
```


```
[35]: # Make predicitions
Y_pred = lr_model.predict(X_test)
```

```
[36]: # Evaluate model
from sklearn.metrics import confusion_matrix, classification_report, roc_curve,
    ↪roc_auc_score

# Confusion matrix
cnf_matrix = confusion_matrix(Y_test, Y_pred)
cnf_matrix
```

```
[36]: array([[13,  2],
          [ 3, 16]], dtype=int64)
```


```
[37]: # Prettier confusion matrix (heatmap)
sns.heatmap(pd.DataFrame(cnf_matrix), annot=True, cmap="YlGnBu",
    ↪xticklabels=["Male", "Female"], yticklabels=["Male", "Female"])
plt.xlabel("Predicted class")
plt.ylabel("Actual class")
plt.show();
```



```
[38]: # Classification report
print(classification_report(Y_test, Y_pred, target_names=["Male", "Female"]))
```

	precision	recall	f1-score	support
Male	0.81	0.87	0.84	15
Female	0.89	0.84	0.86	19
accuracy			0.85	34
macro avg	0.85	0.85	0.85	34
weighted avg	0.86	0.85	0.85	34

```
[39]: # ROC curve
Y_pred_proba = lr_model.predict_proba(X_test)[: , 1]
fpr, tpr, _ = roc_curve(Y_test, Y_pred_proba)
auc = roc_auc_score(Y_test, Y_pred_proba)
plt.plot(fpr, tpr, label=f"R2 score: {auc}")
plt.show();
```



4 Time Series Analysis

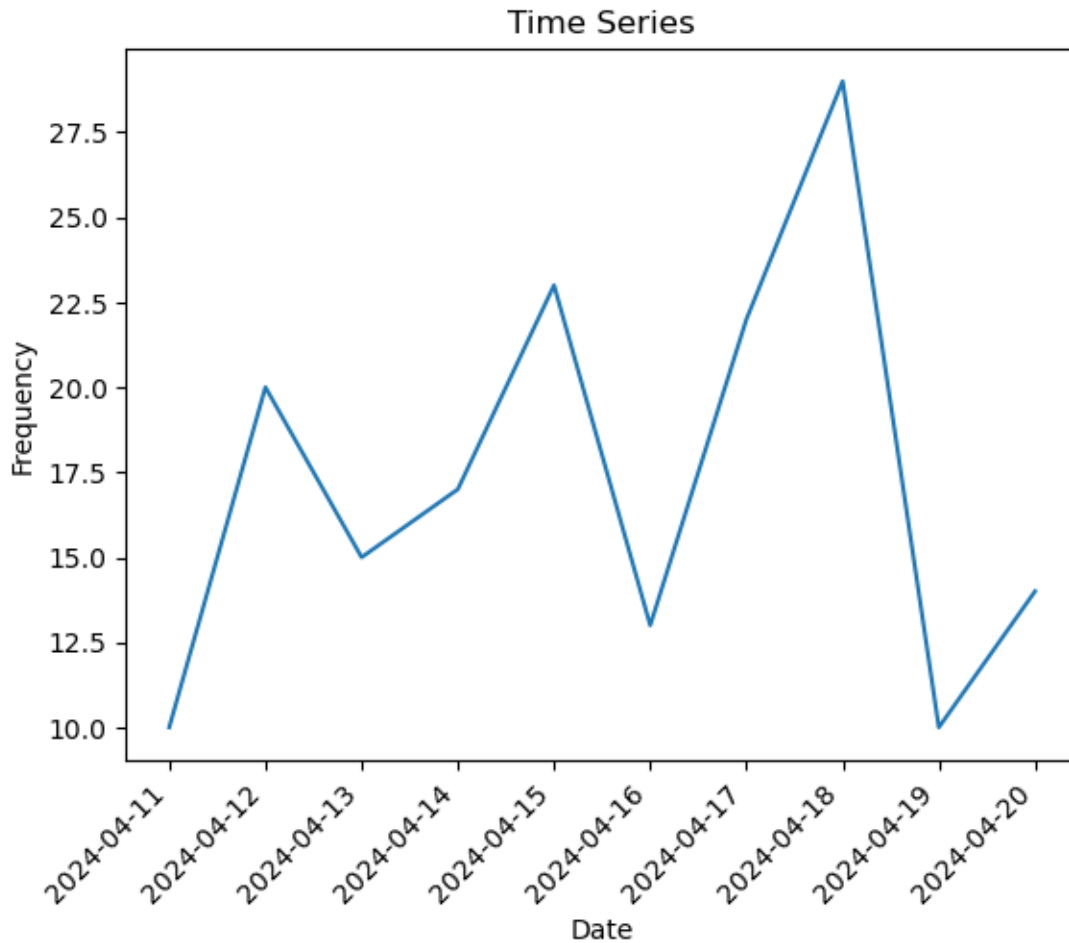
- This example is generating dummy data

```
[40]: # Essential imports
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
```

```
[41]: import datetime

df = pd.DataFrame({'date': np.array([datetime.datetime(2024, 4, i + 1) for i in
↪range(10, 20)]),
                    'frequency': [10, 20, 15, 17, 23, 13, 22, 29, 10, 14]})

plt.plot(df.date, df.frequency)
plt.title("Time Series")
plt.xticks(rotation=45, ha="right")
plt.xlabel("Date")
plt.ylabel("Frequency")
plt.show();
```



5 Data Visualisation (Python)

- This example generates random data for plotting using NumPy
- Of course, you can choose to make your own small sample dataset (recommended)

```
[42]: # Essential libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

```
[43]: # Generate a dummy dataset
np.random.seed(42) # Set random seed for reproducibility
data = pd.DataFrame({
    'A': np.random.rand(100),
    'B': np.random.rand(100),
    'C': np.random.rand(100),
    'Category': np.random.choice(['X', 'Y', 'Z'], 100)
```


```
})

# Display it
data.head()
```


```
[43]:
```

	A	B	C	Category
0	0.374540	0.031429	0.642032	X
1	0.950714	0.636410	0.084140	Y
2	0.731994	0.314356	0.161629	Z
3	0.598658	0.508571	0.898554	Z
4	0.156019	0.907566	0.606429	Z


```
[44]: # Scatter plot
plt.figure(figsize=(8, 6))
plt.scatter(data['A'], data['B'])
plt.title('Scatter plot')
plt.xlabel('A')
plt.ylabel('B')
plt.show();
```



```
[45]: # Line plot
plt.figure(figsize=(8, 6))
plt.plot(data['A'], data['B'])
plt.title('Line plot')
plt.xlabel('A')
plt.ylabel('B')
plt.show();
```



```
[46]: # Histogram
plt.figure(figsize=(8, 6))
plt.hist(data['C'], bins=20)
plt.title('Histogram')
plt.xlabel('C')
plt.ylabel('Frequency')
plt.show();
```

```
[47]: # Box plot
plt.figure(figsize=(8, 6))
data.boxplot(column=['A', 'B', 'C'])
plt.title('Box plot')
plt.ylabel('Values')
plt.show();
```

